

CMPUT 328 Fall 2024

Assignment 1

Worth 8% of the total weight

Part 1: Logistic Regression for MNIST [worth 2% of the total weight]:

Implement Logistic Regression in PyTorch (You can make use of the multiple linear regression notebook from lecture):

- Train and test on MNIST by defining your own data pipeline for training, validation and testing using PyTorch dataloader.
 - Use the last 12,000 samples of the train set as validation set.
 - Test the trained model on the validation set every few epochs to prevent overfitting.
 - Do not use the test set for training.**
 - Use Stochastic Gradient Descent ([SGD](#)) as an optimizer.
 - Use the CrossEntropy loss.
- Add a **regularization term** to improve your model (L1 or L2 regularization, whichever gives better accuracy)

Expected Performance: A correctly implemented, and somewhat well-tuned version of this algorithm will have an accuracy of **92-94%** on both test and validation sets of MNIST.

You need to complete the function *logistic_regression* in *AI_submission.py* for this part. The return type of the function is a dictionary where you have to assign your trained logistic regression model to the variable **model = None**.

Part 2: Fully-connected Neural Network for CIFAR10[Worth 4% of the total weight]

Implement a Fully-connected Neural Network using the built-in functions of PyTorch. Train this network on the CIFAR10 dataset with CrossEntropy loss. The CIFAR10 dataset each image dimension is 32x32x3 as it is a color RGB image.

Please refer to the *Network Architecture* section for the specifications.

You need to complete the class *FNN* in *AI_submission.py* to implement the forward pass as well as the loss computation.

- `__init__(self, loss type, num classes)` initializes your network layers
- `forward(self, x)` takes a batch of images as a tensor of size $N \times (32 \times 32 \times 3)$ and returns the class probabilities as a tensor of size $N \times 10$ where N is the batch size
- `get_loss(self, output, target)` takes the output of the forward pass and ground truth labels of the corresponding images and returns a tensor containing the loss computed according to the loss type argument of `__init__`

Network Architecture

$$Y_p = \text{Softmax}(\text{Relu}(\text{Tanh}(XW_1 + b_1)W_2 + b_2))W_3 + b_3)$$

You can use built-in torch functions for defining the layers (`nn.Linear`) and activations. Dimensions of the vectors and matrices are as follows: X contains the input images having a shape $(N \times (32 \times 32 \times 3))$. N is the batch size. W_1 is $(32 \times 32 \times 3 \times 64)$, b_1 is (1×64) , W_2 is (64×32) , b_2 is (1×32) , W_3 is (32×10) , b_3 is (1×10) . Output probabilities Y_p has the shape $(N \times 10)$. Note that for each of N indices in the first dimension, the softmax function is applied along the second dimension of its input matrix.

Part 3: Hyperparameter Search [Worth 2% of the total weight]

Find optimal hyperparameters using Adaptive Moment Estimation ([Adam](#)) as an optimizer on both part 1 (Logistic Regression) (1% of the part-3 weight) and part 2 (FNN) (1% of the part-3 weight).

- You *should* perform **grid search** or **random search** for finding the optimal hyper-parameters using accuracy on the validation set and select the best configuration.
- You can also use more advanced search strategies like evolutionary search, but you are **not** allowed to use any automatic parameter search methods like *scorch*.
- You **cannot** use the test set during this process.

You need to complete the function *tune_hyper_parameter* in *AI_submission.py* for this part. Since you require the best parameters and best metric for each of the models, keep them as a list of dictionaries, where the first dictionary in the list contains best parameters and metric for part-1 and the second dictionary for part-2. It's up to you how you want to keep the dictionary for each part as long as the parameters and metrics names used as keys in the dictionary are self explanatory.

Template Code

You are provided with template code in the form of three files: *AI_main.py*, *FNN_main.py* and *AI_submission.py*.

You need to complete the two functions (i.e. *logistic_regression* for part 1, *tune_hyper_parameter* for part 3) and *FNN* class in

AI_submission.py. You can add any other functions or classes you want to *AI_submission.py* but do not make any changes to *FNN_main.py* and *AI_main.py*.

Running the code

Your own machine

Install python (version ≥ 3.6) if needed and install the required packages by running:

python3 -m pip install numpy torch torchvision tqdm paramparse

Run the code using:

python3 AI_main.py

python3 FNN_main.py

It is recommended to use an IDE like pycharm or vscode to make debugging easier.

Colab

Run this from a code cell in notebooks:

!python3 "<full path to AI_main.py >"

!python3 "<full path to FNN_main.py >"

You can optionally install the *paramparse* package to enable command line arguments:

!python3 -m pip install paramparse

You can then use command line arguments as:

part 1:

!python3 "<full path to AI_main.py >" mode=logistic

part 2:

!python3 "<full path to FNN_main.py >" mode=fnn

part 3:

!python3 "<full path to FNN_main.py >" mode=tune target_metric=accuracy

Submission

You need to submit only the completed *AI_submission.py*. Make sure to import any additional libraries you need so it can be used as a standalone Python module from *FNN_main.py* and *AI_main.py*.

To reiterate, please **do not** submit *FNN_main.py* and *AI_main.py* or any other files generated by running the code.

Marking

Part 1: Marks will depend on correctness of the implementation along with the following metrics:

- **Runtime:** The total runtime of your submission (including training and testing) **should not exceed 300 seconds** for either dataset on Colab GPU.
 - One trick to improve your run time is to grid search the hyperparameters first but only put in the best hyperparameters you found in your submission.
- **Accuracy:** Score scales linearly from **83 - 93%** accuracy on the test set

Part 2: Marks will depend on correctness of the implementation along with the following metrics:

- **Runtime:** The total runtime of your submission (including training and testing) **should not exceed 300 seconds** for either dataset on Colab GPU.
- **Accuracy:** Score scales linearly from **37-42%** accuracy on the test set

Part 3: Marks will depend on the correctness of your search implementation.

- **Runtime:** The runtime of your submission **should not exceed 1500 seconds** on Colab GPU.
- **Accuracy:** There are no specific accuracy requirements except there should be improvement in loss / accuracy compared to the baseline

Runtime Penalty:

If you exceed the runtime threshold by **10*k%**, you will be penalized **k%**.

For example if you exceed the runtime by 20% then the incurred penalty will be 2%.

